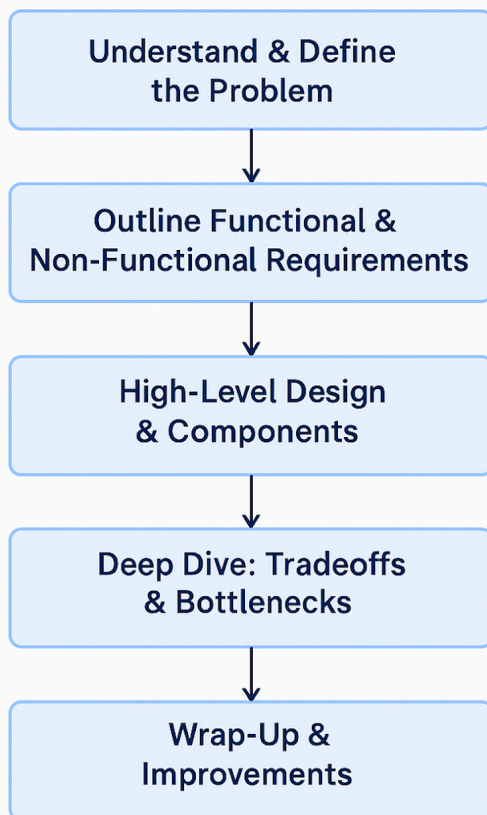


System Design Prep - By Kartikey (Meta | Google | Amazon)

System Design Interview Flow



System Design Cheat Sheet

Real-Time Systems

- Messaging System (WhatsApp): Delivery guarantees, offline support, scalable infra
- Notification System (Push/SMS/Email): Token mgmt, retries, multi-channel delivery
- Ride Sharing (Uber/Lyft): ETA, geo-matchmaking, surge logic
- Uber Matching: Rider-driver matching, surge pricing, location updates
- Webhook Callback: Reliable delivery, retries, exactly-once guarantees

Financial & Commerce Systems

- Payment System (Stripe): Idempotent flows, retry handling, PCI compliance
- Ticket Booking System (BookMyShow): Atomic seat locks, payment consistency, race condition handling
- E-Commerce Platform (Amazon): Catalog, cart, order mgmt, scalability

- Ad Click Aggregator (Google Ads): High-throughput ingestion, deduplication, analytics
- Ad Click Analytics: Windowed aggregation, Kafka ingestion, low-latency queries

Discovery & Social Systems

- NewsFeed System (Instagram): Timeline construction, fanout, ranking
- Trending Hashtags (Twitter): Sliding window top-K, approx. counting, real-time updates
- Search Autocomplete (Google): Trie lookups, ranking, typo-tolerance
- Voting System (Reddit): Idempotent votes, fraud prevention, tallying
- Coding Platform (LeetCode): Problem mgmt, submission flow, eval infra

Infra & Platform Components

- Distributed Cache (Redis/Memcached): Fast R/W, cache invalidation, eviction strategies
 - Rate Limiter (API Gateway): Token/leaky bucket, TTL, distributed enforcement
 - Amazon Lockers (IoT Logistics): Order to locker mapping, real-time status, auth
 - Web Crawler (Googlebot): URL queuing, deduplication, rate limiting
 - Task Scheduler (Cron): Recurring jobs, retries, cron parsing
 - Task Management Tool (Jira): CRUD, background jobs, notifications
 - Log Analysis (ELK Stack): Real-time ingestion, indexing, anomaly detection
 - Image Hosting (Imgur/S3): Scalable storage, deduplication, CDN integration
 - URL Shortener (Bitly): Unique code generation, high-QPS redirect, link analytics
-

Mock Interview Flow + Questions (Deep Dive)

0–5 mins: Understand & Define the Problem

- Clarify Requirements: What does the system do? What's out-of-scope?
- Target Users & Scale: Is it for 1M or 100M users? Concurrency matters.
- Access Patterns: Read-heavy or write-heavy? Real-time vs batch?
- Example: For a ride-sharing app, clarify scope (payments? ratings? live tracking?)

5–15 mins: Functional & Non-Functional Requirements

- Functional Examples (for Uber): Request a ride, Match driver, Track ride, Pay, Rate
- Non-Functional Goals:
 - Latency: <100ms for critical APIs
 - Availability: 99.99%
 - Durability: No message loss
 - Scalability: Handle spikes with horizontal scale
- Capacity Planning:
 - 100M DAUs, 500k QPS, 80:20 read/write ratio

15–30 mins: High-Level Design & Components

- Sketch Architecture:
 - API Gateway → Service Layer → Cache → DB → Queue
 - Modular services: MatchingService, UserService, BillingService
- Database:
 - SQL: PostgreSQL/MySQL (for consistency)
 - NoSQL: MongoDB/Cassandra (for scale/flexibility)
- Caching:
 - Redis for hot reads, TTL strategies, invalidation
- Queues:
 - Kafka/SQS for async flows: billing, notifications
- API Design:
 - RESTful/GraphQL, idempotent POSTs, versioning

30–40 mins: Deep Dive — Tradeoffs, Bottlenecks, Scaling

- Data Consistency Models:
 - Strong (ticketing), Eventual (feeds), Causal (messaging)
- CAP Tradeoffs:
 - Choose CA, CP, or AP based on product goals
- Scaling:
 - Sharding by user/location, leader-follower replication
 - Load balancers with health checks
- Failures:
 - Circuit breakers, retry/backoff, dead letter queues
- Observability:
 - Logging (structured), metrics, alerts

40–45 mins: Wrap-Up & Future Enhancements

- Summarize the full system workflow
- Identify weak links and improvement paths
- Scope V2 features:
 - Multi-region active-active
 - Analytics pipelines
 - GDPR & PCI compliance

Mental Frameworks (Advanced)

Design Thinking

- Always start with user experience and intent
- APIs define contracts, versioning ensures evolution

Scalability & Performance

- Design for 10x scale even if MVP
- Use event-driven async processing where possible
- Latency budgets per component (e.g. 30ms DB, 10ms cache)
- Optimize read-heavy vs write-heavy differently

System Reliability

- Think in terms of failure paths
- Add retries, rate limits, graceful degradation
- Secure every component (encryption, RBAC, audit logs)

Cost + Maintainability

- Monitor cloud usage: storage, compute, transfer
 - Favor managed services if ops is not differentiating
-

7-Day System Design Mastery Plan

Day 1: URL Shortener + Ticket Booking

Goal: ID generation, atomic ops, consistency guarantees

Day 2: Messaging System + Payment Infra

Goal: Delivery guarantees, retries, PCI-compliance

Day 3: Rate Limiter + Cache System

Goal: High QPS handling, token buckets, eviction logic

Day 4: Newsfeed + Voting System

Goal: Timeline construction, ranking, idempotency

Day 5: E-Commerce Platform + Search Autocomplete

Goal: DB schema design, search latency reduction

Day 6: Uber Matching + Notification System

Goal: Geo-matching, mobile token handling

Day 7: Mock Interview + Design Review

Goal: Simulate a full 45-min interview, get feedback

System Design Tradeoffs Summary

Use this reference during interviews or system evaluations to justify architecture decisions.

Tradeoff Area	Option A	Option B
Consistency	Strong (ACID, strict correctness)	Eventual (faster, scalable)
Storage	SQL (structured, transactional)	NoSQL (flexible, scalable)
Communication	Synchronous (instant feedback)	Asynchronous (decoupled, resilient)
Latency Handling	Push (instant updates)	Pull (scalable, client-driven)
Scaling	Vertical (bigger machine)	Horizontal (more machines, sharding)
Data Access	Read-optimized (cache, replicas)	Write-optimized (batch, queues)
State Management	Centralized (single source of truth)	Distributed (high availability)
Queue Processing	At-most-once (fast)	At-least-once / Exactly-once (safe, retry logic)
Caching	Aggressive (low latency)	Minimal (high accuracy)
API Design	Monolith (easy to start)	Microservices (scalable, modular)